

Dimensionality Reduction

Algorithms in Machine Learning, ISAE-SUPAERO

Fabrice Jimenez
Data Scientist & Artificial Intelligence Engineer
Airbus Commercial Aircraft

Dimensionality reduction: why?


What is high dimension?

v1	v2	v3	v4	...	v365
8.4	15	2.2	0.5	...	65.8
9.1	10	5.1	-4.3	...	-7
...


Many, many features... Maybe more features than data points!

Let's consider 10 points in 20 dimensions v1 ... v20


4D?
nD??




First issue: **visualization**



2D



3D



The Curse of Dimensionality

One so-called “Curse”, several effects...

Effect n°1 - Data Coverage

Statistics, Machine Learning: generalization / estimation of population, from a reduced **sample which is representative**


Representative = training set covering “enough” portions of feature space

Imagine we need to train a binary classifier (red or blue points)...




How would you classify the new green point?


1D



2D



3D



The logo for 'The Curse of Monkey Island' features the title in a large, stylized, yellow font with a black outline. Above the title, a small pirate character is visible. To the right of the title, there is a large, glowing treasure chest with a flame-like effect. The background is black.

Ideally, training samples “uniformly” distributed on the feature space...

If in 1D you need ~10 points:


- in 2D you need ~100 points
- in 3D you need ~1000 points
- in nD you need $\sim 10^n$ points

For fixed number of samples, many dimensions = poor coverage
= classifier with poor performance!!!


1D



2D



3D



The Curse of Dimensionality

One so-called “Curse”, several effects...

Effect n°2 - Distance concentration

With high dimension, **sparsity of points increases**, and Euclidean distance becomes more and more equal between all points

Notion of distance important for many ML tasks: clustering, outlier detection...

Many dimensions = distance obsolete = results obsolete...

Let's take our dataset with 10 points in 20 dimensions:

- For each dimension d , keep the d first variables
- Compute all pairwise distances
- Study the relative difference between distances




More details:

On the surprising behavior of distance metrics in high dimensional space

Charu C. Aggarwal et. al.

$$\lim_{d \rightarrow \infty} E \left(\frac{\text{dist}_{\max}(d) - \text{dist}_{\min}(d)}{\text{dist}_{\min}(d)} \right) \rightarrow 0$$



Roughly, when d exceeds the number of points, distance becomes useless...

The Curse of Dimensionality


One so-called “Curse”, several effects...

Effect n°3 - Noise pollution

When looking for a pattern in low dimension, if you add many dimensions irrelevant to this pattern: **noise conceals the actual information**

Relative difference between distances:

- Main jump between 1D and 2D (introduction of noise)
- Second jump roughly when $d > \text{number of points}$ (effect n°2)




Can you identify the **pattern**?


1D



2D



3D



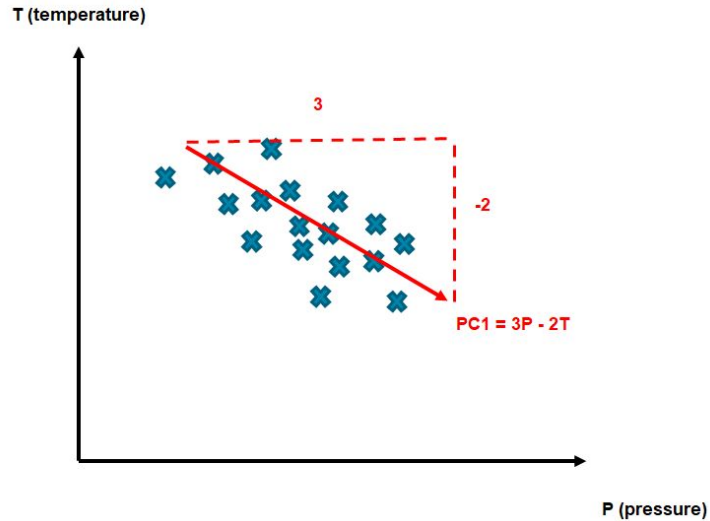
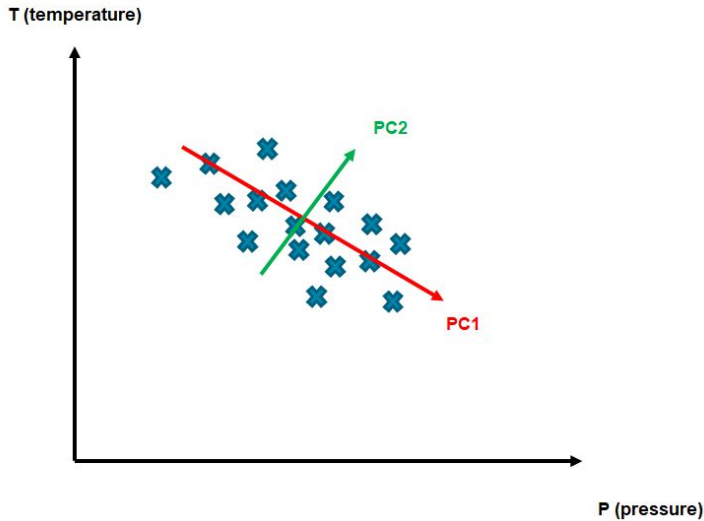
PCA and information pattern

Effects especially valid if dimensions are drawn from independent distributions

Correlations reinforce patterns visibility, but **how to distinguish correlation information, noise, and pattern information ?**

Principal Components Analysis → Linear combination of initial features

- First components maximize the projected variability of the dataset = spread distances = often **main information patterns**
- Correlated features are gathered in the same components: **removes correlation redundancy**



Defeat the Curse of Dimensionality? Compute PCA and keep the first principal components might be a solution...

PCA Computation

Dataset Matrix

$$M = \begin{bmatrix} X_{11} & \dots & X_{1p} \\ \dots & \dots & \dots \\ X_{n1} & \dots & X_{np} \end{bmatrix}$$

Centering

$$\bar{M} = \begin{bmatrix} X_{11} - \bar{X}_1 & \dots & X_{1p} - \bar{X}_p \\ \dots & \dots & \dots \\ X_{n1} - \bar{X}_1 & \dots & X_{np} - \bar{X}_p \end{bmatrix}$$

Scaling

$$\tilde{M} = \begin{bmatrix} \frac{X_{11} - \bar{X}_1}{\sigma(X_1)} & \dots & \frac{X_{1p} - \bar{X}_p}{\sigma(X_p)} \\ \dots & \dots & \dots \\ \frac{X_{n1} - \bar{X}_1}{\sigma(X_1)} & \dots & \frac{X_{np} - \bar{X}_p}{\sigma(X_p)} \end{bmatrix}$$

Covariance matrix: $cov(M) = \frac{1}{n} \times \bar{M}^T \cdot \bar{M}$

Correlation matrix: $cor(M) = \frac{1}{n} \times \tilde{M}^T \cdot \tilde{M}$

Eigenvalue decomposition: $A = V^T \cdot D \cdot V$ where $D = diag(\lambda_1, \lambda_2, \dots, \lambda_n)$

- applies to square matrices $n \times n$ which are diagonalizable
- **V**: columns of **V** are eigenvectors V_i for the corresponding eigenvalues λ_i i.e. $AV_i = \lambda_i V_i$

PCA Computation:

- **Correlation matrix is square, symmetric, real:** diagonalizable in orthonormal basis (Spectral theorem...)
- **Principal Components:** eigenvectors of correlation matrix
- **Explained variance by component i:** i th eigenvalue

Exercise: PCA and diagonalization



Demonstrate that principal components of the PCA are the eigenvectors of the correlation matrix

1/ Consider principal component with direction vector p
Our dataset M projected on this component becomes F such as:


$$F = \tilde{M}p \quad \begin{cases} \bar{F} = 0 \\ \|p\| = 1 \end{cases}$$

2/ Write the explained variance on this component:

$$\begin{aligned} \text{var}(F) &= \frac{1}{n} \sum_{i=1}^n (F_i - \bar{F})^2 = \frac{1}{n} \sum_{i=1}^n F_i^2 = \frac{1}{n} F^T F \\ &= \frac{1}{n} (\tilde{M}p)^T \tilde{M}p \\ &= \frac{1}{n} p^T (\tilde{M}^T \tilde{M}) p \\ \text{var}(F) &= p^T \text{cor}(M) p \end{aligned}$$

4/ Explained variance on the component is then:

$$\text{var}(F) = p^T \text{cor}(M) p = p^T \mu p = \mu \quad \leftarrow \text{Explained variance is eigenvalue of the correlation matrix}$$



3/ We want the component to explain a maximum of variance, we need to formulate the following maximization problem:


$$\begin{aligned} \begin{cases} \text{maximize } [p^T \text{cor}(M) p] \\ \|p\| = p^T p = 1 \end{cases} &\iff \begin{cases} \mathcal{L}(p) = p^T \text{cor}(M) p - \mu(p^T p - 1) \\ \frac{\partial \mathcal{L}}{\partial p}(p) = 2\text{cor}(M)p - 2\mu p = 0 \end{cases} \\ &\iff \text{cor}(M)p = \mu p \end{aligned}$$

p must be eigenvector of the correlation matrix

PCA Interpretation

Example 1:

One of the variables “draws” all the variance of the dataset: careful with scaling!



No scaling: 1 variable with high variance “draws” all PCA effect to itself

Scaling: 1 noise variable will have same variance as information variable


If different units: scaling mandatory, otherwise does not make sense!

Be careful with quick interpretations!


Example 2:

Be careful with proximity of points in projection graph

Green and blue points
very close...



Actually no!
But not visible in
the projection
PC1 - PC2

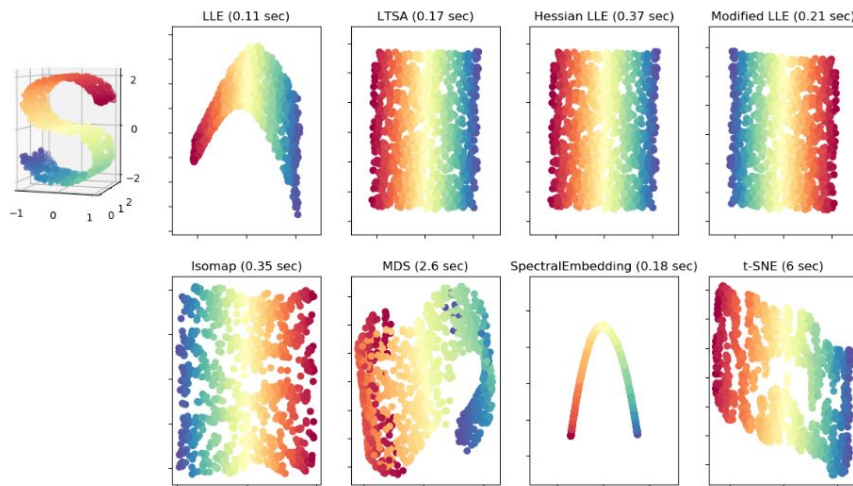


Think about these examples when you get your results...

PCA vs Manifold Learning

PCA is nice when target pattern is linked to variance in a linear direction...

Unfortunately, it's not always the case: patterns may be non-linear. Non-linear projection methods exist to reduce dimension:



(source: scikit-learn)

Manifold Learning

We will see 1 example in this course: t-SNE (t-distributed Stochastic Neighbor Embedding)

t-distributed Stochastic Neighbor Embedding (t-SNE)


Particular technique of dimensionality reduction: it transcribes the similarities between points in a lower dimension

Similarities in initial space: normal probability density around each point

More details on t-SNE:

Visualizing data using t-SNE

Laurens van der Maaten et. al.



Perplexity: parameter
ruling the coverage
→ to change sigma

$$P_{ij} = \frac{\exp(-||x_i - x_j||^2 / 2\sigma_i^2)}{\sum_{k \neq i} \exp(-||x_i - x_k||^2 / 2\sigma_i^2)}$$

then

$$P_{ij} = \frac{P_{ij} + P_{ji}}{2n}$$


t-distributed Stochastic Neighbor Embedding (t-SNE)

Particular technique of dimensionality reduction: it transcribes the similarities between points in a lower dimension

Similarities in reduced space: Cauchy probability density around each point

$$Q_{ij} = \frac{(1 + \|y_i - y_j\|^2)^{-1}}{\sum_{k \neq i} (1 + \|y_i - y_k\|^2)^{-1}}$$

Learning rate and stopping criteria:
parameters impacting
accuracy of the projection



Objective:

minimize difference between distance distributions in initial and reduced space

How:

optimization (gradient descent or others) on the Kullback-Leibler divergence $KL(P, Q)$

Going from normal to Cauchy distribution: stretches the distances → highly differentiated clusters (sometimes too much!)

Pros: high transcription capability of complex patterns

Cons: sensitive to tuning (optimization algorithm)

Dimensionality Reduction: Application

It's time to play on your own...

Main interest = discover the dataset, play with PCA, t-SNE parameters, and take a step back for interpretation



Questions?

